

COMPSCI 689

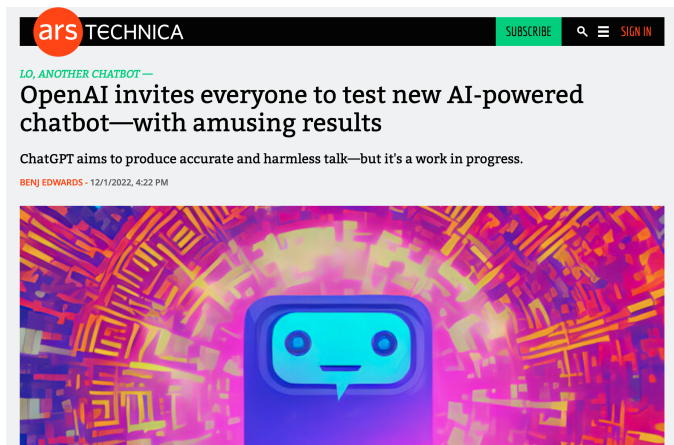
Lecture 21: Transformers

Benjamin M. Marlin

College of Information and Computer Sciences
University of Massachusetts Amherst

Slides by Benjamin M. Marlin (marlin@cs.umass.edu).

ChatGPT - 12/1/2022



<https://arstechnica.com/information-technology/2022/12/openai-invites-everyone-to-test-new-ai-powered-chatbot-with-amusing-results/>

ChatGPT - 10/01/2024

nature

Explore content ▾

About the journal ▾

Publish with us ▾

Subscribe

[nature](#) > [news](#) > article

NEWS | 01 October 2024

‘In awe’: scientists impressed by latest ChatGPT model o1

The chatbot excels at science, beating PhD scholars on a hard science test. But it might ‘hallucinate’ more than its predecessors.

By [Nicola Jones](#)



<https://www.nature.com/articles/d41586-024-03169-9>

Transformers




Image Credit: Ashish Vaswani et al.

Attention

- The fundamental building block of transformers is a computation referred to as *attention*.
- Given an input consisting of a collection of M objects represented as column vectors $\mathbf{x}_i$ of length D , the attention computation outputs a new representation of the M objects as column vectors $\mathbf{h}_i$ of length K .
- The attention computation allows for arbitrary mixing of information across the collection of objects regardless of the number of objects.

Computing Attention

- The self attention computation first transforms each input vector $\mathbf{x}_i$ into three different embeddings referred to as the key $\mathbf{k}_i$, the query $\mathbf{q}_i$ and the value $\mathbf{v}_i$, each of length K .
- All the transformations are linear. The parameters W^k , W^q and W^v are all $K \times D$ matrices.

$$\mathbf{k}_i = W^k \cdot \mathbf{x}_i$$

$$\mathbf{q}_i = W^q \cdot \mathbf{x}_i$$

$$\mathbf{v}_i = W^v \cdot \mathbf{x}_i$$

Computing Attention

- The next step computes an attention weight a_{ij} for each pair of input objects (i, j) .
- a_{ij} indicates how much object i will pay attention to object j when computing the output representation $\mathbf{h}_i$.
- The attention weights for object i are the softmax transform of the similarity between the query $\mathbf{q}_i$ for object i and the keys $\mathbf{k}_j$ for each object j in the collection, with scaling applied:

$$a_{ij} = \frac{\exp(\frac{1}{\sqrt{k}} \mathbf{q}_i^T \mathbf{k}_j)}{\sum_{j'=1}^T \exp(\frac{1}{\sqrt{k}} \mathbf{q}_i^T \mathbf{k}_{j'})}$$

Computing Attention

- The final output representation $\mathbf{h}_i$ for each object i is obtained as an attention weighted linear combination of the values $\mathbf{v}_j$ of all objects:

$$\mathbf{h}_i = \sum_{j=1}^K a_{ij} \mathbf{v}_j$$

- We can use this computation to define a self attention block that provides a mapping from a $M \times D$ representation of the M objects $\mathbf{x}$ to an $M \times K$ representation $\mathbf{h}$:

$$\mathbf{h} = \text{SelfAttention}_{\theta}(\mathbf{x})$$

Computing Self-Attention

- The parameters θ of the block self-attention $_{\theta}(\mathbf{x})$ are the three weight matrices $\theta = [W_k, W_q, W_v]$. The total parameter count is thus $3DK$.
- Note that similar to an RNN and CNN and unlike an MLP, the number of parameters used in the self-attention block is independent of the number of objects. It does depend on the length D of the individual input vectors.

Multi-Head Attention

- A single self-attention computation is often referred to as an “attention head.”
- We can make models more complex by learning L different attention heads with different parameters θ_l .
- The multi-head attention computation can also produce a final representation of vectors of length k using an additional combination of the outputs of the multiple attention heads:

$$\mathbf{h}'_l = \text{SelfAttention}_{\theta_l}(\mathbf{x})$$

$$\mathbf{h}_i = \sum_{l=1}^L W_l^c \mathbf{h}'_{li}$$

Multi-Head Attention

- We can use the multi-head attention computation to define a multi-head attention block that also provides a mapping from a $M \times D$ representation of M objects $\mathbf{x}$ to an $M \times K$ representation $\mathbf{h}$:

$$\mathbf{h} = \text{MultiheadAttention}_{\phi}(\mathbf{x})$$

- The parameters ϕ include the parameters θ_l for each of the L self-attention blocks, as well as the multi-head attention combination weights W^c .

Transformers




Image Credit: Ashish Vaswani et al.

Transformer Blocks

- A transformer block is a multi-head attention-based module that also performs a mapping from a collection of M objects in D dimensions to a collection of M objects in K dimensions.
- It starts with the multi-head attention computation and then performs several additional layers of non-linear transformations that act independently on the representations output by the multi-head attention process:

$$\mathbf{h} = \text{MultiheadAttention}_{\phi}(\mathbf{x})$$

$$\mathbf{u}_i = \text{LayerNorm}(\mathbf{x}_i + \mathbf{h}_i, \gamma_1, \beta_1)$$

$$\mathbf{z}'_i = W_2 \text{relu}(W_1 \mathbf{u}_i)$$

$$\mathbf{z}_i = \text{LayerNorm}(\mathbf{u}_i + \mathbf{z}'_i, \gamma_2, \beta_2)$$

Aside: Layer Norm

- The LayerNorm block performs a re-scaling of the values within each vector:

$$\text{LayerNorm}(\mathbf{z}, \gamma, \beta) = \beta + \gamma \frac{(\mathbf{z} - \mu(\mathbf{z}))}{\sigma(\mathbf{z})}$$

$$\mu(\mathbf{z}) = \frac{1}{K} \sum_{k=1}^K \mathbf{z}_k$$

$$\sigma(\mathbf{z}) = \left(\frac{1}{K} \sum_{k=1}^K (\mathbf{z}_k - \mu(\mathbf{z}))^2 \right)^{1/2}$$

Transformer Blocks

- We can use the transformer block computation to define a new transformation computation:

$$\mathbf{z} = \text{TransformerBlock}_{\omega}(\mathbf{x})$$

- Here the parameters ω include the parameters ϕ of the included multi-head attention computation, and the parameters $\gamma_1, \gamma_2, \beta_1, \beta_2, W_1$ and W_2 used in the rest of the transformer block computation.

Transformers

- A transformer is simply a stack of R transformer blocks each with their own parameters ω_r .
- Each block in the stack takes it's input from the previous block, except for the first block, which takes it's input from $\mathbf{x}$.

$$\mathbf{z}^1 = \text{TransformerBlock}_{\omega_1}(\mathbf{x})$$

$$\mathbf{z}^2 = \text{TransformerBlock}_{\omega_2}(\mathbf{z}^1)$$

$$\vdots$$

$$\mathbf{z}^R = \text{TransformerBlock}_{\omega_R}(\mathbf{z}^{R-1})$$

Transformer Applications

- Since the transformer operates on representations of a collection of objects, it is a highly general model.
- The original applications of transformers were in natural language processing. They were then extended to models for multivariate time series and images.

Transformers For Sequences

- For sequential data such as text and time series, the first step in applying a transformer is to form a multi-dimensional representation of each sequence element.
- For multivariate time series, the data at each time point is already a vector of some length D .
- For text, it is common to embed the data at the word level. This can be done using fixed one-hot encoding, a fixed previously learned embedding, or a learnable embedding.

Representing Structure

- Since the transformer treats its inputs as a collection of objects, it is actually invariant to the sequential ordering of the inputs.
- As a result, transformers need to be provided with additional data about the structure of the collection of objects.
- For time series and text, the model must be provided with a representation of the sequential ordering of the objects. This is generally referred to as a “positional encoding.”

Positional Encodings

- A simple way to provide a transformer block with positional information is to use a one-hot encoding $p(i)$ of the position of each input data vector $\mathbf{x}_i$ in the sequence.
- This one-hot vector can be concatenated with the original data vector $\tilde{\mathbf{x}}_i = [\mathbf{x}_i; p(i)]$. The transformer then operates using $\tilde{\mathbf{x}}_i$ as input.
- Note that when stacking transformer blocks, we may want to re-inject positional encoding information at the input to each block in the stack.

Positional Encodings

- In the original transformers paper, the proposed model uses a different approach that can be thought of as producing an embedding of positional information $p(i)$ of size D for each position i in the input sequence.
- Instead of concatenating the position information with the feature information, the original model adds the positional encoding vector to the data vectors yielding $\tilde{\mathbf{x}}_i = \mathbf{x}_i + p(i)$. The transformer then operates using $\tilde{\mathbf{x}}_i$ as input.
- Again, this can be repeated at the input to each transformer block.

Transformers




Image Credit: Ashish Vaswani et al.

Models for Text

- In the case of text data, current transformer models are trained as unsupervised models by learning to predict the next word in a sequence.
- The output of the final transformer block is a $M \times K$ matrix.
- We can select a given context window size C and use the final representations in this window to predict the probability of the next word in the sequence.

$$p(W_{M+1} = w | \mathbf{x}_1, \dots, \mathbf{x}_M) = \frac{\exp \left(\sum_{i=1}^C \sum_{k=1}^K W_{wik}^p \cdot \mathbf{z}_{ik}^R \right)}{\sum_{w'=1}^V \exp \left(\sum_{i=1}^C \sum_{k=1}^K W_{w'ik}^p \cdot \mathbf{z}_{ik}^R \right)}$$

Text Transformer Model Parameters




GPT4 is reported to have 1.5 trillion parameters. Note: the human brain has about 85 billion neurons and 100 trillion synapses.

GPT Models Sizes

Model Name	n_{params}	n_{layers}	d_{model}	n_{heads}	d_{head}	Batch Size	Learning Rate
GPT-3 Small	125M	12	768	12	64	0.5M	6.0×10^{-4}
GPT-3 Medium	350M	24	1024	16	64	0.5M	3.0×10^{-4}
GPT-3 Large	760M	24	1536	16	96	0.5M	2.5×10^{-4}
GPT-3 XL	1.3B	24	2048	24	128	1M [*]	2.0×10^{-4}
GPT-3 2.7B	2.7B	32	2560	32	80	1M	1.6×10^{-4}
GPT-3 6.7B	6.7B	32	4096	32	128	2M	1.2×10^{-4}
GPT-3 13B	13.0B	40	5140	40	128	2M	1.0×10^{-4}
GPT-3 175B or “GPT-3”	175.0B	96	12288	96	128	3.2M	0.6×10^{-4}

RLHF



ChatGPT Errors

ChatGPT ▾

How many r's are in strawberry?




The word "strawberry" contains one letter "r."

Where are the r's?



In the word "strawberry," the letter "r" appears twice:

1. The first "r" is in the 4th position: **strawberry**.
2. The second "r" is in the 7th position: **strawberry**.



ChatGPT Ethics

90.5 WESA
Public Radio for Western Piedmont

NEWSLETTERS

SIGN IN

NPR SHOP

DONATE

NEWS CULTURE MUSIC PODCASTS & SHOWS SEARCH



BOOKS

LISTEN & FOLLOW



Some authors are suing OpenAI. Will it backfire?

NOVEMBER 16, 2023 · 4:28 PM ET

HEARD ON [ALL THINGS CONSIDERED](#)By [Keith Romer](#)

3-Minute Listen

+ PLAYLIST



Fiction writers like George R.R. Martin and Jonathan Franzen are suing OpenAI for using their books to train ChatGPT. That lawsuit could paradoxically benefit the company being sued.

<https://www.npr.org/2023/11/16/1213588978/>

LLMs and Sustainability

≡

SEARCH

FORTUNE


SIGN IN

Subscribe Now

SUCCESS · CLIMATE CHANGE

AI doesn't just require tons of electric power. It also guzzles enormous sums of water.

BY JANE THIER
September 23, 2024 at 2:44 PM EDT



In order to shoot off one email per week for a year, ChatGPT would use up 27 liters of water, or about one-and-a-half jugs.

JASMIN MERDAD - GETTY IMAGES

[https://fortune.com/2024/09/23/
ai-water-usage-droughts-chatgpt-microsoft/](https://fortune.com/2024/09/23/ai-water-usage-droughts-chatgpt-microsoft/)

LLMs and Sustainability



World ▾ US Election Business ▾ Markets ▾ More ▾

My New

Google to buy power for AI needs from small modular nuclear reactor company Kairos

By Timothy Gardner

October 14, 2024 3:46 PM EDT · Updated a month ago



LLM Risks



[Featured](#) [Topics](#) [Newsletters](#) [Events](#) [Podcasts](#)

[SIGN IN](#)

[SUBSCRIBE](#)

ARTIFICIAL INTELLIGENCE

Geoffrey Hinton tells us why he's now scared of the tech he helped build

"I have suddenly switched my views on whether these things are going to be more intelligent than us."

By Will Douglas Heaven

May 2, 2023

<https://www.technologyreview.com/2023/05/02/1072528/>